

autism_detection

April 8, 2025

1 Importing required Libraries and the Dataset.

```
[6]: import numpy as np
import pandas as pd
from time import time
from IPython.display import display # Allows use of Display() for Dataframes.

# Import supplementary visualization code visuals.py
import visuals as vs

%matplotlib inline

data = pd.read_csv("autism_data.csv")
display(data.head(5))
```

	A1_Score	A2_Score	A3_Score	A4_Score	A5_Score	A6_Score	A7_Score	\
0	1	1	1	1	0	0	1	
1	1	1	0	1	0	0	0	
2	1	1	0	1	1	0	1	
3	1	1	0	1	0	0	1	
4	1	0	0	0	0	0	0	

	A8_Score	A9_Score	A10_Score	...	gender	ethnicity	jundice	austim	\
0	1	0	0	...	f	White-European	no	no	
1	1	0	1	...	m	Latino	no	yes	
2	1	1	1	...	m	Latino	yes	yes	
3	1	0	1	...	f	White-European	no	yes	
4	1	0	0	...	f	?	no	no	

	contry_of_res	used_app_before	result	age_desc	relation	Class/ASD
0	United States		no	6.0 18 and more	Self	NO
1	Brazil		no	5.0 18 and more	Self	NO
2	Spain		no	8.0 18 and more	Parent	YES
3	United States		no	6.0 18 and more	Self	NO
4	Egypt		no	2.0 18 and more	?	NO

[5 rows x 21 columns]

```
[7]: # Total number of records:
n_records = len(data.index)

# Total number of records with ASD
n_asd_yes = len(data[data['Class/ASD'] == 'YES'])

# Total number of records without ASD
n_asd_no = len(data[data['Class/ASD'] == 'NO'])

# Percentage of individuals with ASD
yes_percentage = float((n_asd_yes) / n_records * 100)

# Printing the outputs
print(f'Total number of records : {n_records}')
print(f'Number of individuals with ASD : {n_asd_yes}')
print(f'Number of individuals without ASD : {n_asd_no}')
print("Percentage of individuals with ASD : {:.2f}%".format(yes_percentage))
```

```
Total number of records : 704
Number of individuals with ASD : 189
Number of individuals without ASD : 515
Percentage of individuals with ASD : 26.85%
```

1.1 Featureset Exploration

```
[8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 704 entries, 0 to 703
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   A1_Score               704 non-null   int64
1   A2_Score               704 non-null   int64
2   A3_Score               704 non-null   int64
3   A4_Score               704 non-null   int64
4   A5_Score               704 non-null   int64
5   A6_Score               704 non-null   int64
6   A7_Score               704 non-null   int64
7   A8_Score               704 non-null   int64
8   A9_Score               704 non-null   int64
9   A10_Score              704 non-null   int64
10  age                    702 non-null   float64
11  gender                 704 non-null   object
12  ethnicity               704 non-null   object
13  jundice                 704 non-null   object
14  austim                 704 non-null   object
15  contry_of_res          704 non-null   object
```

```

16  used_app_before  704 non-null    object
17  result           704 non-null    float64
18  age_desc         704 non-null    object
19  relation          704 non-null    object
20  Class/ASD         704 non-null    object
dtypes: float64(2), int64(10), object(9)
memory usage: 115.6+ KB

```

```
[9]: data.describe()
```

```

[9]:
      count  A1_Score  A2_Score  A3_Score  A4_Score  A5_Score  A6_Score  \
count  704.000000  704.000000  704.000000  704.000000  704.000000  704.000000
mean    0.721591    0.453125    0.457386    0.495739    0.498580    0.284091
std     0.448535    0.498152    0.498535    0.500337    0.500353    0.451301
min     0.000000    0.000000    0.000000    0.000000    0.000000    0.000000
25%     0.000000    0.000000    0.000000    0.000000    0.000000    0.000000
50%     1.000000    0.000000    0.000000    0.000000    0.000000    0.000000
75%     1.000000    1.000000    1.000000    1.000000    1.000000    1.000000
max     1.000000    1.000000    1.000000    1.000000    1.000000    1.000000

      count  A7_Score  A8_Score  A9_Score  A10_Score  age  result
count  704.000000  704.000000  704.000000  704.000000  702.000000  704.000000
mean    0.417614    0.649148    0.323864    0.573864    29.698006    4.875000
std     0.493516    0.477576    0.468281    0.494866    16.507465    2.501493
min     0.000000    0.000000    0.000000    0.000000    17.000000    0.000000
25%     0.000000    0.000000    0.000000    0.000000    21.000000    3.000000
50%     0.000000    1.000000    0.000000    1.000000    27.000000    4.000000
75%     1.000000    1.000000    1.000000    1.000000    35.000000    7.000000
max     1.000000    1.000000    1.000000    1.000000   383.000000   10.000000

```

1.2 Preparing the Data

```
[10]: data.isna().sum()
```

```

[10]: A1_Score      0
      A2_Score      0
      A3_Score      0
      A4_Score      0
      A5_Score      0
      A6_Score      0
      A7_Score      0
      A8_Score      0
      A9_Score      0
      A10_Score     0
      age           2
      gender         0
      ethnicity      0

```

```
jundice      0
austim       0
contry_of_res 0
used_app_before 0
result       0
age_desc     0
relation     0
Class/ASD    0
dtype: int64
```

Dropping missing values

```
[11]: data.dropna(inplace=True)
data.describe()
```

```
[11]:
```

	A1_Score	A2_Score	A3_Score	A4_Score	A5_Score	A6_Score \
count	702.000000	702.000000	702.000000	702.000000	702.000000	702.000000
mean	0.723647	0.452991	0.458689	0.497151	0.498575	0.284900
std	0.447512	0.498140	0.498646	0.500348	0.500354	0.451689
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	A7_Score	A8_Score	A9_Score	A10_Score	age	result
count	702.000000	702.000000	702.000000	702.000000	702.000000	702.000000
mean	0.417379	0.650997	0.324786	0.574074	29.698006	4.883191
std	0.493478	0.476995	0.468629	0.494835	16.507465	2.498051
min	0.000000	0.000000	0.000000	0.000000	17.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	21.000000	3.000000
50%	0.000000	1.000000	0.000000	1.000000	27.000000	4.000000
75%	1.000000	1.000000	1.000000	1.000000	35.000000	7.000000
max	1.000000	1.000000	1.000000	1.000000	383.000000	10.000000

```
[12]: # After Data Cleaning

# Total number of records:
n_records = len(data.index)

# Total number of records with ASD
n_asd_yes = len(data[data['Class/ASD'] == 'YES'])

# Total number of records without ASD
n_asd_no = len(data[data['Class/ASD'] == 'NO'])

# Printing the outputs
print("AFTER REMOVING NULL VALUES : ")
```

```
print(f'Total number of records : {n_records}')
print(f'Number of individuals with ASD : {n_asd_yes}')
print(f'Number of individuals without ASD : {n_asd_no}')
```

AFTER REMOVING NULL VALUES :

Total number of records : 702


Number of individuals with ASD : 189

Number of individuals without ASD : 513


1.2.1 Visualizations with Seaborn

```
[13]: import seaborn as sns
import matplotlib.pyplot as plt
sns.set(style="whitegrid", color_codes=True)
```


```
[14]: # Draw a nested violinplot and split the violins for easier comparison
sns.violinplot(x="result", y="jundice", hue="austim", data=data,
               split=True, inner="quart", palette={'yes': "r", 'no': "b"})
sns.despine(left=True)
```



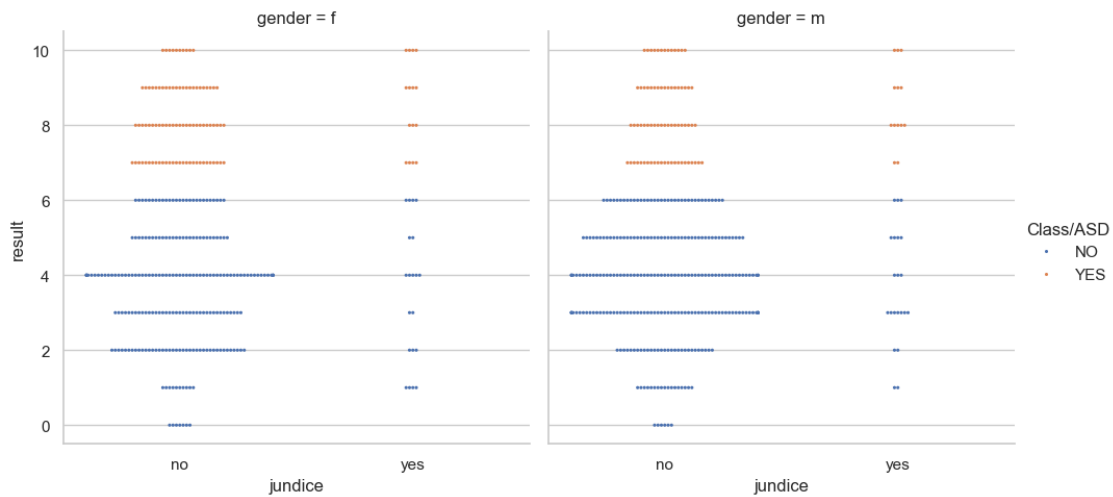
```
[15]: # Draw a nested violinplot and split the violins for easier comparison
sns.violinplot(x="result", y="jundice", hue="Class/ASD", data=data, split=True,
              inner="quart", palette={'YES': "r", 'NO': "b"})
sns.despine(left=True)
```



```
[16]: # Draw a nested violinplot and split the violins for easier comparison
sns.violinplot(x="result", y="gender", hue="Class/ASD", data=data, split=True,
              inner="quart", palette={'YES': "r", 'NO': "b"})
sns.despine(left=True)
```



```
[17]: sns.catplot(x="jundice", y="result", hue="Class/ASD", s = 5, col="gender",
    ↪ data=data, kind="swarm");
```



Convert the Pandas dataframes into numpy arrays that can be used by scikit_learn. Let's create

an array that extracts only the feature data we want to work with and another array that contains the classes (class/ASD).

```
[18]: data_raw = data['Class/ASD']
features_raw = data[['age', 'gender', 'ethnicity', 'jundice', 'austim',
↪ 'contry_of_res', 'result',
↪
↪ 'relation', 'A1_Score', 'A2_Score', 'A3_Score', 'A4_Score', 'A5_Score', 'A6_Score', 'A7_Score', 'A8_
↪ 'A9_Score', 'A10_Score']]
```

Data Preprocessing : using MinMaxScaler()

```
[19]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
num = ['age', 'result']
features_minmax_transform = pd.DataFrame(data = features_raw)
features_minmax_transform[num] = scaler.fit_transform(features_raw[num])
```

```
[20]: display(features_minmax_transform.head(5))
```

	age	gender	ethnicity	jundice	austim	contry_of_res	result	\
0	0.024590	f	White-European	no	no	United States	0.6	
1	0.019126	m	Latino	no	yes	Brazil	0.5	
2	0.027322	m	Latino	yes	yes	Spain	0.8	
3	0.049180	f	White-European	no	yes	United States	0.6	
4	0.062842	f	?	no	no	Egypt	0.2	

	relation	A1_Score	A2_Score	A3_Score	A4_Score	A5_Score	A6_Score	\
0	Self	1	1	1	1	0	0	
1	Self	1	1	0	1	0	0	
2	Parent	1	1	0	1	1	0	
3	Self	1	1	0	1	0	0	
4	?	1	0	0	0	0	0	

	A7_Score	A8_Score	A9_Score	A10_Score
0	1	1	0	0
1	0	1	0	1
2	1	1	1	1
3	1	1	0	1
4	0	1	0	0

1.2.2 One-Hot Encoding on features_minmax_transform

```
[21]: features_final = pd.get_dummies(features_minmax_transform)
features_final.head(5)
```

```
[21]:      age  result  A1_Score  A2_Score  A3_Score  A4_Score  A5_Score  \
0  0.024590    0.6         1         1         1         1         0
```


1	0.019126	0.5	1	1	0	1	0
2	0.027322	0.8	1	1	0	1	1
3	0.049180	0.6	1	1	0	1	0
4	0.062842	0.2	1	0	0	0	0

	A6_Score	A7_Score	A8_Score	...	contry_of_res_United Kingdom	\
0	0	1	1	...		False
1	0	0	1	...		False
2	0	1	1	...		False
3	0	1	1	...		False
4	0	0	1	...		False

	contry_of_res_United States	contry_of_res_Uruguay	contry_of_res_Viet Nam	\
0	True	False		False
1	False	False		False
2	False	False		False
3	True	False		False
4	False	False		False

	relation_?	relation_Health care professional	relation_Others	\
0	False		False	False
1	False		False	False
2	False		False	False
3	False		False	False
4	True		False	False

	relation_Parent	relation_Relative	relation_Self
0	False	False	True
1	False	False	True
2	True	False	False
3	False	False	True
4	False	False	False

[5 rows x 103 columns]

1.2.3 Encode all classes data to numerical values

```
[22]: data_classes = data_raw.apply(lambda x : 1 if x == 'YES' else 0)
```

```
[23]: encoded = list(features_final.columns)
print("{} total features after one-hot encoding".format(len(encoded)))
print(encoded)
```

```
103 total features after one-hot encoding
['age', 'result', 'A1_Score', 'A2_Score', 'A3_Score', 'A4_Score', 'A5_Score',
'A6_Score', 'A7_Score', 'A8_Score', 'A9_Score', 'A10_Score', 'gender_f',
'gender_m', 'ethnicity_?', 'ethnicity_Asian', 'ethnicity_Black',
'ethnicity_Hispanic', 'ethnicity_Latino', 'ethnicity_Middle Eastern ',
```

```

'ethnicity_Others', 'ethnicity_Pasifika', 'ethnicity_South Asian',
'ethnicity_Turkish', 'ethnicity_White-European', 'ethnicity_others',
'jundice_no', 'jundice_yes', 'austim_no', 'austim_yes',
'contry_of_res_Afghanistan', 'contry_of_res_AmericanSamoa',
'contry_of_res_Angola', 'contry_of_res_Argentina', 'contry_of_res_Armenia',
'contry_of_res_Aruba', 'contry_of_res_Australia', 'contry_of_res_Austria',
'contry_of_res_Azerbaijan', 'contry_of_res_Bahamas', 'contry_of_res_Bangladesh',
'contry_of_res_Belgium', 'contry_of_res_Bolivia', 'contry_of_res_Brazil',
'contry_of_res_Burundi', 'contry_of_res_Canada', 'contry_of_res_Chile',
'contry_of_res_China', 'contry_of_res_Costa Rica', 'contry_of_res_Cyprus',
'contry_of_res_Czech Republic', 'contry_of_res_Ecuador', 'contry_of_res_Egypt',
'contry_of_res_Ethiopia', 'contry_of_res_Finland', 'contry_of_res_France',
'contry_of_res_Germany', 'contry_of_res_Hong Kong', 'contry_of_res_Iceland',
'contry_of_res_India', 'contry_of_res_Indonesia', 'contry_of_res_Iran',
'contry_of_res_Iraq', 'contry_of_res_Ireland', 'contry_of_res_Italy',
'contry_of_res_Japan', 'contry_of_res_Jordan', 'contry_of_res_Kazakhstan',
'contry_of_res_Lebanon', 'contry_of_res_Malaysia', 'contry_of_res_Mexico',
'contry_of_res_Nepal', 'contry_of_res_Netherlands', 'contry_of_res_New Zealand',
'contry_of_res_Nicaragua', 'contry_of_res_Niger', 'contry_of_res_Oman',
'contry_of_res_Pakistan', 'contry_of_res_Philippines', 'contry_of_res_Portugal',
'contry_of_res_Romania', 'contry_of_res_Russia', 'contry_of_res_Saudi Arabia',
'contry_of_res_Serbia', 'contry_of_res_Sierra Leone', 'contry_of_res_South
Africa', 'contry_of_res_Spain', 'contry_of_res_Sri Lanka',
'contry_of_res_Sweden', 'contry_of_res_Tonga', 'contry_of_res_Turkey',
'contry_of_res_Ukraine', 'contry_of_res_United Arab Emirates',
'contry_of_res_United Kingdom', 'contry_of_res_United States',
'contry_of_res_Uruguay', 'contry_of_res_Viet Nam', 'relation_?',
'relation_Health care professional', 'relation_Others', 'relation_Parent',
'relation_Relative', 'relation_Self']

```

```

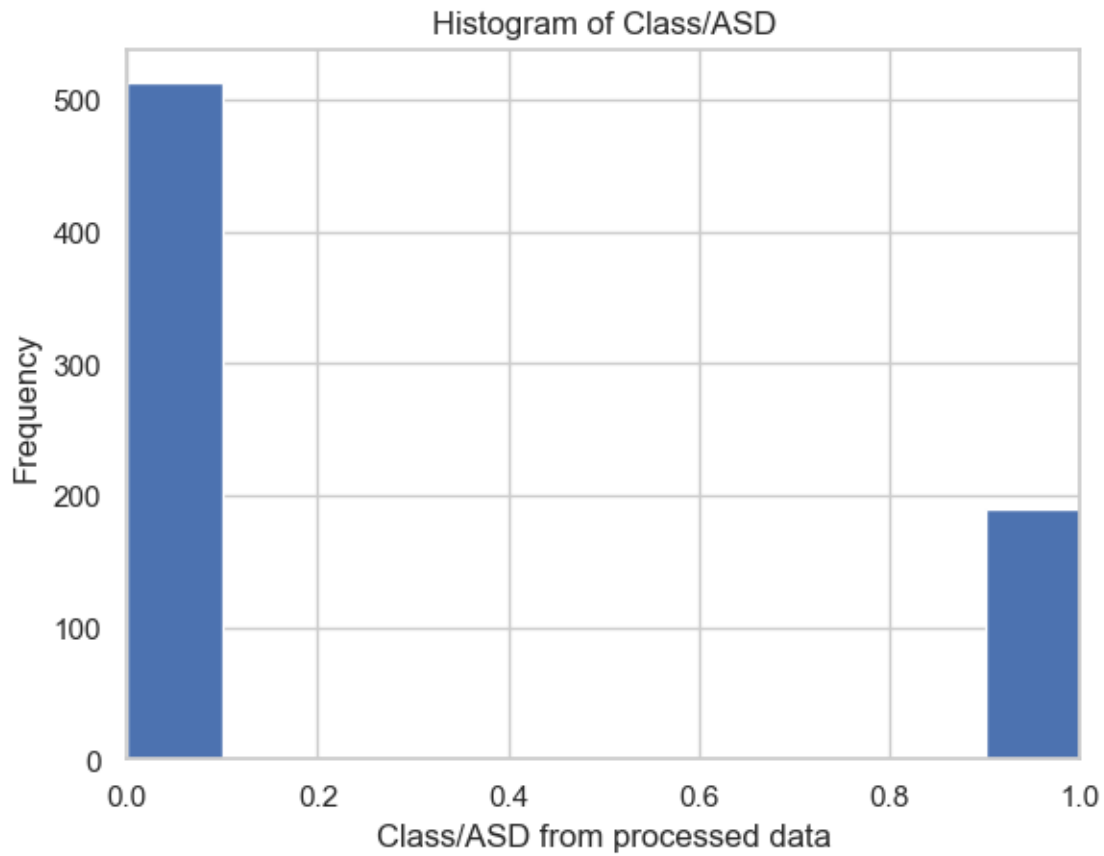
[24]: plt.hist(data_classes, bins=10)
      plt.xlim(0,1)
      plt.title('Histogram of Class/ASD')
      plt.xlabel('Class/ASD from processed data')
      plt.ylabel('Frequency')

```

```

[24]: Text(0, 0.5, 'Frequency')

```



1.3 Shuffle and Split the data

All the categorical variables have been converted to numerical variables and have been normalized, we split the data into train and test set, test set will be 20% of the total data.

```
[25]: from sklearn.model_selection import train_test_split
      np.random.seed(123)
      X_train, X_test, y_train, y_test = \
          train_test_split(features_final, data_classes, test_size=0.2, random_state=1)
      print("Train set has {} enteries.".format(X_train.shape[0]))
      print("Test set has {} enteries.".format(X_test.shape[0]))
```

Train set has 561 enteries.

Test set has 141 enteries.

1.4 Models :

1.4.1 1. Decision Tress

```
[26]: from sklearn import tree
      from sklearn.tree import DecisionTreeClassifier
      dec_model = DecisionTreeClassifier()
      dec_model.fit(X_train.values, y_train)
```

```
[26]: DecisionTreeClassifier()
```

```
[27]: y_pred = dec_model.predict(X_test.values)
      print('True : ', y_test.values[0:25])
      print('False : ', y_pred[0:25])
```

```
True :  [1 0 0 0 0 0 1 1 0 0 1 1 0 0 0 0 0 0 0 0 0 1 1 0]
False :  [1 0 0 0 0 0 1 1 0 0 1 1 0 0 0 0 0 0 0 0 0 1 1 0]
```

```
[28]: from sklearn import metrics
      cm = metrics.confusion_matrix(y_test, y_pred)
      print(cm)
      TP = cm[1,1]
      FP = cm[0,1]
      TN = cm[0,0]
      FN = cm[1,0]
```

```
[[101   0]
 [  0  40]]
```

```
[29]: print('Accuracy:')
      print((TN+TP)/float(TP+TN+FP+FN))
      print('Error:')
      print((FP+FN)/float(TP+TN+FP+FN))
      print('Precision:')
      print(metrics.precision_score(y_test,y_pred))
      print('Score:')
      print(dec_model.score(X_test.values, y_test))
```

```
Accuracy:
1.0
Error:
0.0
Precision:
1.0
Score:
1.0
```

1.4.2 2. Random Forest

```
[30]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
rndm_model = RandomForestClassifier(n_estimators=5, random_state=1)
cv_score = cross_val_score(rndm_model, features_final, data_classes, cv =10)
cv_score.mean()
```

```
[30]: np.float64(0.9900603621730383)
```

```
[31]: # F-beta Score
rndm_model.fit(X_train.values, y_train)
from sklearn.metrics import fbeta_score
y_pred = rndm_model.predict(X_test.values)
fbeta_score(y_test, y_pred, average='binary', beta=0.5)
```

```
[31]: 1.0
```

1.4.3 3. Support Vector Machine

```
[32]: from sklearn import svm
svm_model = svm.SVC(kernel='linear', C=1, gamma=2)
cv_score = cross_val_score(svm_model, features_final, data_classes, cv =10)
cv_score.mean()
```

```
[32]: np.float64(1.0)
```

```
[33]: #F-beta Score
svm_model.fit(X_train.values, y_train)
from sklearn.metrics import fbeta_score
y_pred = svm_model.predict(X_test.values)
fbeta_score(y_test, y_pred, average='binary', beta=0.5)
```

```
[33]: 1.0
```

1.4.4 4. K-Nearest-Neighbors(KNN)

```
[34]: from sklearn import neighbors
knn_model = neighbors.KNeighborsClassifier(n_neighbors=10)
cv_score = cross_val_score(knn_model, features_final, data_classes, cv =10)
cv_score.mean()
```

```
[34]: np.float64(0.9458752515090543)
```

```
[35]: #F-beta Score
knn_model.fit(X_train.values, y_train)
from sklearn.metrics import fbeta_score
y_pred = knn_model.predict(X_test.values)
```

```
fbeta_score(y_test, y_pred, average='binary', beta=0.5)
```

```
[35]: 0.9183673469387755
```

```
[36]: for n in range(10,30):  
      knn_model = neighbors.KNeighborsClassifier(n_neighbors=n)  
      cv_scores = cross_val_score(knn_model, features_final, data_classes, cv=10)  
      print (n, cv_scores.mean())
```

```
10 0.9458752515090543  
11 0.9473239436619719  
12 0.9444869215291751  
13 0.9501609657947686  
14 0.9458953722334005  
15 0.9458953722334004  
16 0.951569416498994  
17 0.951549295774648  
18 0.9529778672032194  
19 0.9572635814889336  
20 0.9529778672032194  
21 0.9529778672032194  
22 0.9486921529175051  
23 0.9472635814889336  
24 0.9486921529175051  
25 0.9486720321931589  
26 0.9515090543259557  
27 0.9501006036217303  
28 0.9486720321931589  
29 0.9472434607645874
```

Hence, K is not making any significant difference on accuracy of our predictions.

1.4.5 5. Naive Bayes

```
[37]: from sklearn.naive_bayes import MultinomialNB  
      nb_model = MultinomialNB()  
      cv_score = cross_val_score(nb_model, features_final, data_classes, cv =10)  
      cv_score.mean()
```

```
[37]: np.float64(0.8746277665995976)
```

```
[38]: #F-beta Score  
      nb_model.fit(X_train.values, y_train)  
      from sklearn.metrics import fbeta_score  
      y_pred = nb_model.predict(X_test.values)  
      fbeta_score(y_test, y_pred, average='binary', beta=0.5)
```

```
[38]: 0.7675438596491229
```

1.4.6 6. Logistic Regression

```
[39]: from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression()
cv_score = cross_val_score(lr_model, features_final, data_classes, cv =10)
cv_score.mean()
```

```
[39]: np.float64(0.9971428571428571)
```

```
[40]: #F-beta Score
lr_model.fit(X_train.values, y_train)
from sklearn.metrics import fbeta_score
y_pred = lr_model.predict(X_test.values)
fbeta_score(y_test, y_pred, average='binary', beta=0.5)
```

```
[40]: 0.9948979591836735
```

1.5 Model Tuning

```
[43]: from sklearn.metrics import fbeta_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import make_scorer
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.svm import SVC
```

```
[44]: def f_beta_score(y_true, y_predict):
    return fbeta_score(y_true, y_predict, beta = 0.5)
clf = SVC(random_state = 1)
parameters = {'C':range(1,6), 'kernel':
    ↳['linear', 'poly', 'rbf', 'sigmoid'], 'degree':range(1,6)}
scorer = make_scorer(f_beta_score)
```

```
[45]: grid_obj = GridSearchCV(estimator = clf, param_grid = parameters, scoring =
    ↳scorer)
grid_fit = grid_obj.fit(X_train.values, y_train)
best_clf = grid_fit.best_estimator_
```

```
[46]: predictions = (clf.fit(X_train.values, y_train)).predict(X_test.values)
best_predictions = best_clf.predict(X_test.values)
```

```
[47]: print ("Unoptimized model\n-----")
print ("Accuracy score on testing data: {:.4f}".format(accuracy_score(y_test,
    ↳predictions)))
print ("F-score on testing data: {:.4f}".format(fbeta_score(y_test,
    ↳predictions, beta = 0.5)))
print ("\nOptimized Model\n-----")
```

```
print ("Final accuracy score on the testing data: {:.4f}".  
      ↪format(accuracy_score(y_test, best_predictions)))  
print ("Final accuracy score on the testing data: {:.4f}".  
      ↪format(accuracy_score(y_test, best_predictions)))  
print ("Final F-score on the testing data: {:.4f}".format(fbeta_score(y_test, ↪  
      ↪best_predictions, beta = 0.5)))
```

Unoptimized model

Accuracy score on testing data: 0.9645

F-score on testing data: 0.9574

Optimized Model

Final accuracy score on the testing data: 1.0000

Final accuracy score on the testing data: 1.0000

Final F-score on the testing data: 1.0000

[]: